

## C++11 has Thread Support

<http://en.cppreference.com/w/cpp/thread>

However, C++11 (2011!) support is still experimental in g++ 4.8.5 (2015; in use in current stable Linux distros still 2017; gcc 7 (2019) has full C++14 support) and has to be manually turned on using `-std=c++11`:

```
/usr/include/c++/4.8/bits/c++0x_warning.h:32:2:  
error: #error This file requires compiler and library support  
for the ISO C++ 2011 standard. This support is currently experimental,  
and must be enabled with the -std=c++11 or -std=gnu++11 compiler options.
```

The GNU implementation of C++11 threads uses the pthreads library, and therefore the pthreads-library has to be linked!

- The design is completely different from OpenMP, e.g., does not have a simple parallel loop.
- The task-based paradigm is designed after Pthreads (and Boost threads).
- The STL is still not thread safe :(
- However, parallel reads are O.K. in STL containers; parallel writes of different objects O.K.;
- Write conflicts (race conditions) need something like a “critical” in OpenMP;  
in C++11 this functionality is provided by the `mutex` (mutually exclusive) and `lock` class;

en.cppreference.com/w/cpp/thread 240% Suchen

besucht Webmail TUBAF Lehre TUBAF Webmail UzK Login - OTRS Owncloud OTRS-Alle suse package search VideoLAN - VLC media ... BSI BSI Grundschrift neu Liv Rebecca Arnedatter...

**cppreference.com** Create account Search

Page Discussion View Edit History

C++ Thread support library

## Thread support library

C++ includes built-in support for threads, mutual exclusion, condition variables, and futures.

### Threads

Threads enable programs to execute across several processor cores.

Defined in header `<thread>`

<b>thread</b> (C++11)	manages a separate thread (class)
-----------------------	--------------------------------------

### Functions managing the current thread

Defined in namespace `this_thread`

<b>yield</b> (C++11)	suggests that the implementation reschedule execution of threads (function)
----------------------	--

<b>get_id</b> (C++11)	returns the thread id of the current thread (function)
-----------------------	---

<b>sleep_for</b> (C++11)	stops the execution of the current thread for a specified time duration (function)
--------------------------	---

<b>sleep_until</b> (C++11)	stops the execution of the current thread until a specified time point (function)
----------------------------	--

### Cache size access

Defined in header `<new>`

<b>hardware_destructive_interference_size</b> (C++17)	min offset to avoid false sharing
<b>hardware_constructive_interference_size</b> (constant)	max offset to promote true sharing

## Simple Example

```
#include <iostream>
#include <vector>
#include <thread>
// needs: g++ -std=c++11 -pthread

using namespace std;
void testme(int N)
{
    cout<<"Thread: "<<this_thread::get_id()<<"  N ist "<<N<<endl;
}

int main()
{
    thread t1(testme,2); // launches an additional thread which executes testme
    thread t2(testme,3); // and another one
    thread t3(testme,4); // and another one
    testme(1)           // the master thread can of course also do something

    t1.join(); t2.join(); t3.join(); // joins threads back to the master (=OpenMP)
}
```

## Output:

```
> g++ -g -std=c++11 -pthread  cpp14_test_threads.cpp  
> ./a.out
```

```
Thread: 139925282436928  N ist 1
```

```
Thread: 139925265450752  N ist 2
```

```
Thread: 139925257058048  N ist 3
```

```
Thread: 139925248665344  N ist 4
```

“-std=c++11” switches on C++11 support

“-pthread” links pthread lib

## Another Example

```
#include <iostream>
#include <thread>
#include <chrono>
#include <vector>
// needs: g++ -std=c++11 -pthread
using namespace std;

void multiply(const int N, const vector<vector<double> > &a,
             const vector<vector<double> > &b, vector<vector<double> > &c) {
    int i,j,k;
    cout<<"Multiply Thread: "<<this_thread::get_id()<<endl;
    for (i = 0; i < N; i++) {
        for (j = 0; j < N; j++) {
            c[i][j] = 0;
            for (k = 0; k < N; k++) {
                c.at(i).at(j) += a.at(i).at(k) * b.at(k).at(j);
            }
        }
    }
}
```

```
int main()
{
    int i,j,k;

    vector<vector<double> > a(N);
    vector<vector<double> > b(N);
    vector<vector<double> > c(N);

    for(int i=0; i<N; i++) {
        a.at(i).resize(N);
        b.at(i).resize(N);
        c.at(i).resize(N);
    }

    for (i = 0; i < N; i++) {          // fill matrices
        for (j = 0; j < N; j++) {
            a.at(i).at(j) = i+j+1.0;
            b.at(i).at(j) = 1.0/(i+j+1.0);
            c.at(i).at(j) = 0.0;
        }
    }
}
```

```
vector<vector<double> > a1(a); // copy data
vector<vector<double> > b1(b);
vector<vector<double> > c1(c);

vector<vector<double> > a2(a); // copy data
vector<vector<double> > b2(b);
vector<vector<double> > c2(c);

{
    testme(0);
    thread t1(testme, 1); // spawn a test thread
    t1.join();           // join in back to master
}
```

```
{  
    // now spawn two parallel threads  
    // we use different matrices for every thread here!  
  
    thread t1(multiply, N,std::ref(a1),std::ref(b1),std::ref(c1));  
    thread t2(multiply, N,std::ref(a2),std::ref(b2),std::ref(c2));  
  
    //multiply(N,a,b,c); // master thread can also do work in parallel  
  
    t1.join(); t2.join();  
}
```

### Output:

```
Thread: 140281508857664  N ist 0  
Thread: 140281491871488  N ist 1  
Multiply Thread: 140281483478784  
Multiply Thread: 140281491871488
```



## Launching a Group of Threads

```
#include <iostream>
#include <thread>
static const int num_threads = 10;
void func(int tid) {
    std::cout << "Called by thread " << tid << std::endl;
}
void main() {
    std::thread t[num_threads]; //An array of threads
    // std::vector<std::thread> t(num_threads); // or a vector :)

    //Launch a group of threads
    for (int i = 0; i < num_threads; ++i) {
        t[i] = std::thread(func, i);
    }
    std::cout << "Launched from the main\n";
    //Join the threads with the main thread
    for (int i = 0; i < num_threads; ++i) {
        t[i].join();
    }
}
```

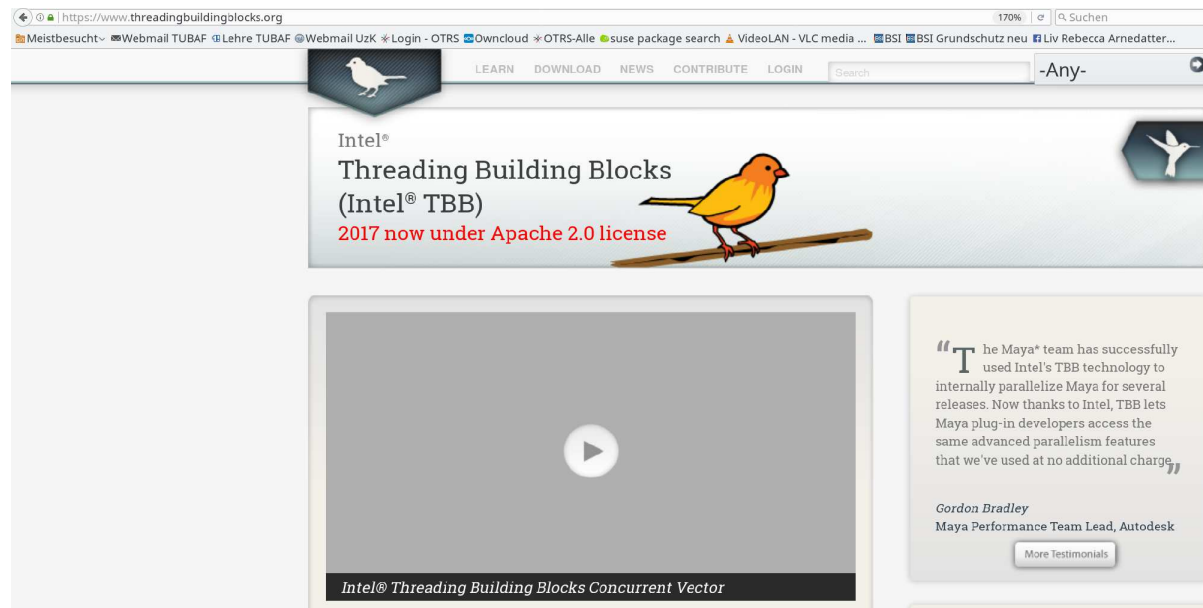
## Example using std::async from cppreference.com

```
#include <iostream>
#include <vector>
#include <algorithm>
#include <numeric>
#include <future>
template <typename RAIter>
int parallel_sum(RAIter beg, RAIter end) {
    auto len = end - beg;
    if(len < 1000)
        return std::accumulate(beg, end, 0);
    RAIter mid = beg + len/2;
    auto handle = std::async(std::launch::async,
                           parallel_sum<RAIter>, mid, end);
    int sum = parallel_sum(beg, mid);
    return sum + handle.get();
}
int main() {
    std::vector<int> v(10000, 1);
    std::cout << "The sum is " << parallel_sum(v.begin(), v.end()) << '\n';
}
```

## For serious parallel algorithms Intel TBB is still much better

TBB = Threads Building Blocks

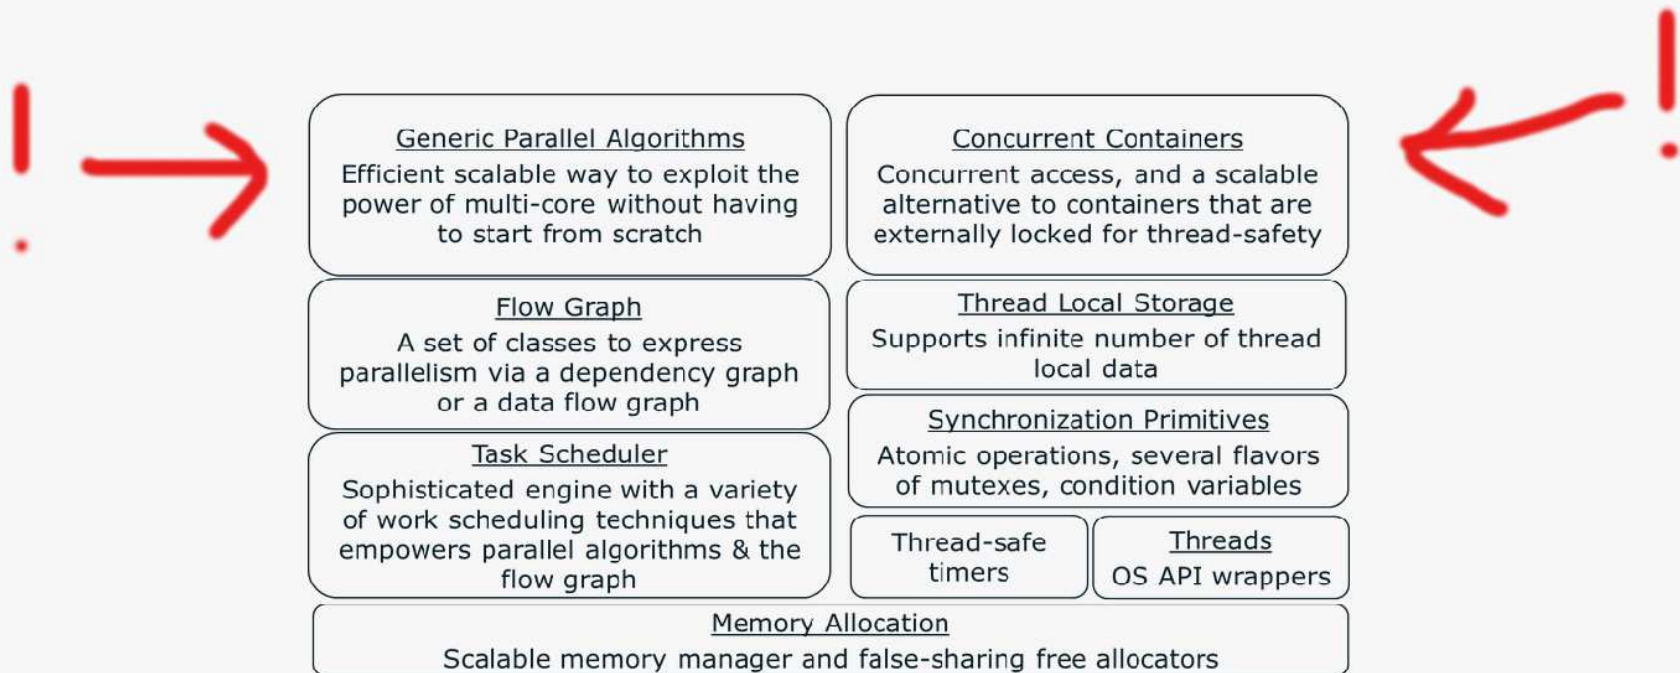
- First version introduced by Intel Corp. in 2006 (!) but is now open (since 2017 under Apache license).
- Threads appeared in the C++ standard only in C++11 (2011)!
- Official Website: <https://www.threadingbuildingblocks.org/>
- Provides serious support for parallel algorithms (unlike C++14)!
- The following is from <https://www.threadingbuildingblocks.org/intel-tbb-tutorial>



## TBB has what OpenMP programmers miss in C++14

### Learn about the Intel® Threading Building Blocks library

The Intel® Threading Building Blocks (Intel® TBB) library provides software developers with a solution for enabling parallelism in C++ applications and libraries. The well-known advantage of the Intel TBB library is that it makes parallel performance and scalability easily accessible to software developers writing loop- and task-based applications. The library includes a number of generic parallel algorithms, concurrent containers, support for dependency and data flow graphs, thread local storage, a work-stealing task scheduler for task based programming, synchronization primitives, a scalable memory allocator and more.



## **TBB has what OpenMP programmers miss in C++14**

Has a parallel for loops (like OpenMP) and (more general!) parallel do loops:

```
#include "tbb/tbb.h"
using namespace tbb;

void ParallelApplyFoo(float a[], size_t n)
{
    tbb::parallel_for(size_t(0), n, [&](size_t i) { Foo(a[i]); } );
}

// example uses Lambda function from C++11
// http://de.cppreference.com/w/cpp/language/lambda
// (but can also be used without)
```

## Excursion: C++11 Lambda Functions

Similar to anonymous functions in Matlab/Octave:

```
% Matlab/Octave:
```

```
>> func=@(x) x+4;
```

```
>> func(6)
```

```
ans = 10
```

```
// TBB + C++11:
```

```
#include <vector>
```

```
#include <iostream>
```

```
#include <algorithm>
```

```
#include <functional>
```

```
int main() // needs -std=c++11
```

```
{
```

```
    std::function<int (int)> func = [](int i) { return i+4; };
```

```
    std::cout << "func: " << func(6) << '\n';
```

```
}
```

## Excursion: C++11 Lambda Functions (2)

Especially usefull as an argument to generic algorithms (here `remove_if`).

```
#include <vector>
#include <iostream>
#include <algorithm>
#include <functional>
int main()  // needs -std=c++11
{
    std::vector<int> c {1,2,3,4,5,6,7};
    int x=5;
    c.erase(std::remove_if(c.begin(),c.end(),[x](int n){ return n < x; } ),c.end());
    std::cout << "c: ";
    for(auto i: c) {
        std::cout << i << ' ';
    }
    std::cout << '\n';
}
```

Therefore also used in TBB for generic parallel algorithms.

## TBB has thread save containers!! - unlike C++14

```
#include <vector>
#include "tbb/tbb.h"
int main(){
    std::vector<int> v; //STL vector - not thread safe, e.g., for size change
    tbb::concurrent_vector<int> w; //TBB vector - thread safe even for push_back()!
}
```

Intel guarantees (<https://software.intel.com/en-us/node/506203>):

*A concurrent\_vector is a container with the following features:*

- *Random access by index. The index of the first element is zero.*
- *Multiple threads can grow the container and append new elements concurrently.*
- *Growing the container does not invalidate existing iterators or indices.*

and “Unlike a `std::vector`, a `concurrent_vector` **never moves existing elements when it grows**. The container **allocates a series of contiguous arrays**. The first reservation, growth, or assignment operation determines the size of the first array. Using a small number of elements as initial size incurs fragmentation across cache lines that may increase element access time. The method `shrink_to_fit()` merges several smaller arrays into a single contiguous array, which may improve access time.”



## **TBB has thread save containers!! - unlike C++14 (2)**

- Much more convenient than STL containers for parallel algorithms.
- A TBB `concurrent_vector` will not move any existing element when the `concurrent_vector` grows: all references and pointers will stay valid.
- An STL `vector` will move existing elements  $\log_2(n)$  times during growth of the vector; adds overhead during growth but no overhead when accessing using `operator []`.
- Therefore, STL vectors are considered as fast as C/C++ builtin arrays and can (=should) replace C/C++ builtin arrays routinely everywhere.
- Of course the TBB strategy will add overhead for the internal management of the memory blocks for every access using `operator []`
- Also, unlike in a `std::vector`, the iterators are not raw pointers.
- For parallel algorithms the benefits of TBB `concurrent_vector` are much more important: It is better to compute the correct result than computing an incorrect result (due to a race condition) fast!